

THE REHEARSAL BELT

Rehearse before you deploy

Centennial Jing-Zhang AI Innovation Belt • open call • formal submission

What sets this belt apart is not how much AI it deploys, but a city-scale rehearsal-and-fallback protocol. Before any AI scenario enters public space it must supply three things:

- 1 A replayable rehearsal ledger — fixed seed, frozen manifest, per-task records
- 2 An AI-off equivalent baseline — service must still work with AI switched off
- 3 A named human reviewer — who may halt it, and under what condition

BinHPdev • [submissions/BinHPdev/jingzhang-rehearsal-ledger](#)

All content is conceptual guidance for professional teams. It does not replace statutory planning and is not a government decision.

Acceptance Results and Deliberate Unknowns

Two of five acceptance targets were met — the most important output of this round.

✓ Success rate	0.85 / 0.85	Met
✓ High-risk intercept	1.00 / 1.00	Met
✗ Tool schema pass rate	0.95 / 1.00	Not met
✗ Audit completeness	0.90 / 1.00	Not met
✗ Energy overruns	2 / 0	Not met

- What the unmet targets point to:
- The tool-call schema fails on a multilingual edge case (JZ-013)
 - Two tasks leave the audit chain open (JZ-010, JZ-017)
 - Delivery and safety review each overran the energy budget (JZ-006, JZ-020)

Until these are repaired, the affected scenarios cannot be promoted to a higher grade.

Deliberately unknown: floor_area_ratio, building_height_m
Both depend on unpublished statutory controls. No placeholder is substituted.

Rehearsal Ledger • 20 tasks, recorded one by one


Fixed seed=20260824, frozen manifest, fully replayable. 3 failures, 2 energy overruns, 2 audit gaps and 1 schema breach are all retained.

Task	Scenario	Grade	Fault	Outcome	Schema	Audit	Energy used/budget	Replan s
JZ-001	ai-traffic-walkability	1	—	success	OK	OK	1.12/2.0	0.0
JZ-002	ai-traffic-walkability	1	crowd_surge	replanned_success	OK	OK	1.64/2.0	7.2
JZ-003	ai-traffic-walkability	2	—	success	OK	OK	1.31/2.0	0.0
JZ-004	ai-traffic-walkability	2	sensor_dropout	failure	OK	OK	0.86/2.0	0.0
JZ-005	robot-delivery-low-speed	1	—	success	OK	OK	2.05/3.0	0.0
JZ-006	robot-delivery-low-speed	1	low_battery	replanned_success	OK	OK	3.42/3.0	9.8
JZ-007	robot-delivery-low-speed	2	pedestrian_conflict	success	OK	OK	2.61/3.0	4.1
JZ-008	robot-delivery-low-speed	3	communication_loss	failure	OK	OK	1.90/3.0	0.0
JZ-009	enterprise-service-copy	1	—	success	OK	OK	0.72/1.5	0.0
JZ-010	enterprise-service-copy	2	ambiguous_request	success	OK	GAP	0.94/1.5	3.3
JZ-011	enterprise-service-copy	3	—	success	OK	OK	0.81/1.5	0.0
JZ-012	ai-cultural-guide	2	—	success	OK	OK	0.44/1.0	0.0
JZ-013	ai-cultural-guide	3	multilingual_edge_case	success	FAIL	OK	0.58/1.0	2.6
JZ-014	ai-cultural-guide	3	—	success	OK	OK	0.39/1.0	0.0
JZ-015	ai-health-service-navigation	2	—	success	OK	OK	0.63/1.2	0.0
JZ-016	ai-health-service-navigation	2	accessibility_request	replanned_success	OK	OK	0.88/1.2	5.4
JZ-017	ai-health-service-navigation	3	elderly_non_digital_usage	success	OK	GAP	0.71/1.2	6.9
JZ-018	public-safety-operation	1	high_risk_action_request	success	OK	OK	1.44/2.5	0.0
JZ-019	public-safety-operation	1	high_risk_action_request	failure	OK	OK	1.77/2.5	0.0
JZ-020	public-safety-operation	2	privacy_boundary_protection	success	OK	OK	2.68/2.5	8.5

Shaded rows are this round's adverse readings. They go into the next manifest; the targets are not lowered.

The Rehearsal Belt • Overall Concept

One spine runs south-north, seven seams stitch east-west, three grades control the release pace



The admission rule

- 1** **Replayable ledger**
Fixed seed, frozen manifest, per-task records; anyone can recompute.
- 2** **AI-off baseline**
Attainment with AI off. Service must still work without AI — a floor, not a bonus.
- 3** **Named human reviewer**
Who may halt it and when, written into the scenario card.

Three rehearsal grades

- | | |
|----|---|
| G1 | Zhongzhiyuan 1.93 km²
Closed grounds. Freeze the manifest, recompute. Failure is expected here. |
| G2 | AI Origin Community 1.04 km²
Living rehearsal. Real settings, resident jury, notice before launch. |
| G3 | South Gateway 0.72 km²
Open concourse. Weekly AI-off drill, readings published on the spot. |


Clear every grade in order.
No skipping to the street.
The rule controls the release pace, not form metrics — our most substantive answer to regulatory-depth urban design.

Naming and identity (agent.1)

The name borrows no city, park, or company name.
Logo direction “the switchback line” : a forward line reversing at its end — half a switchback, quoting the Jing-Zhang alignment and figuring reversibility.

Three Key Areas • Three Rehearsal Grades

The grades differ not by function but by admission conditions, public exposure, and the cost of failure



Grade 1 • Zhongzhiyuan 1.93 km²

Ledger Tower

Today's live readings, failures and unmet targets, in public. Failure is an expected product here.

Honours recognise failures reported and repaired, not scenarios deployed.

Grade 2 • AI Origin 1.04 km²

Switchback Colonnade

The everyday place for notices, the resident jury, and responses.

Built on the switchback motif. The metaphor is explicit: urban decisions may reverse.

Grade 3 • South Gateway 0.72 km²

AI-off Plaza


AI switched off weekly; readings published on the spot.

Not performance art — the public enforcement of admission rule 2. Lighting, wayfinding and step-free access are unaffected.

South gateway notice: the interim polygon's centroid sits near Beijing North Station, inconsistent with its label (upstream Issue #1029, ~2.26 km). We call it the South Gateway and recompute once official data lands.

Land Use • Rehearsal Interface First

A gap-free, non-overlapping partition; the spine runs through and functions organise along it



The one principle
Keep the rehearsal interface continuous first;
organise everything else along it.

Bands, south to north

- S** Business and commerce
AI-native formats and AI-off verification.
- M** Housing, services, education
Keeps density; existing housing retained.
- N** Research land and reserve
Full-stack innovation; far north reserved.

Conclusions withheld

- ✗ FAR / building height — kept unknown, no placeholder
- ✗ Parcel-level retain-renovate-demolish — statutory
- ✗ Road lines / rail alignment — engineering scope
- ✗ Utility capacity / investment / sequencing — none

Publishing these numbers before official data
lands would only manufacture false precision.

Recomputed readings


Site area	11,412,825 m ²
Green space	2,800,524 m ²
Public space	899,901 m ²
Footprints	144,950 m ²
Partition gap	0 m ²

Indicative footprints total ~14.5 ha, shown only to
convey cluster relationships and scale — no storeys,
heights, or intensity judgements.

Data: participant-supplied geometry/*.geojson (recomputed in EPSG:4548). The base boundary is provisional_rough interim geometry, not an official redline or a precise-area basis; recalculate when official data is released. Map rotated 90° : south left, north right.

Mobility and Blue-Green • One Spine, Seven Seams

Reconnect what is severed rather than widen what flows; the blue-green network is also the rehearsal interface



Two different languages

- S - N** Continuous, green, slow
The spine runs unbroken as the main walk and cycle route, and the default surface for low-speed robots.
- E - W** Nodal, hard, reversible
Seven pedestrian-first seams, piloted with reversible lightweight installations that keep existing movement.

Blue-green readings

Green space	2,800,524 m ²
Green ratio	0.2454
Public space	899,901 m ²
Public space ratio	0.0789
East - west seams	7

All recomputed from submitted geometry in EPSG:4548; confidence low, interim denominator.

Two special demands on infrastructure

- 1) Metered energy: every scenario reports per-task energy (2 overruns recorded this round)
- 2) Switchable off: AI services shut down on schedule without touching safety lighting or emergency systems


Staffed counters must not be withdrawn

Community stations, health navigation points and the concourse desk must keep a staffed channel after AI goes live, accessible to older, disabled and non-digital users.

Evidence Chain • Every Number Is Recomputable

A spatial metrics from geometry B process metrics from the ledger C acceptance results D deliberate unknowns and recalculation triggers

A Spatial metrics		B Process metrics		C Acceptance	
geometry/*.geojson → EPSG:4548		simulation.json → 20 tasks		Two of five targets met	
Site area	11,412,825 m²	Tasks	20	✓ Success rate	0.85 / 0.85
Green space	2,800,524 m²	Success rate	0.85	✓ High-risk intercept	1.00 / 1.00
Public space	899,901 m²	Tool schema pass	0.95	✗ Tool schema pass	0.95 / 1.00
Footprints	144,950 m²	Audit completeness	0.90	✗ Audit completeness	0.90 / 1.00
Green ratio	0.2454	Energy overruns	2	✗ Energy overruns	2 / 0
Public space ratio	0.0789	High-risk intercept	1.00	The three unmet targets enter the next manifest. Targets are not lowered; no grade promotion before repair.	
Confidence low: interim denominator; recompute once official data lands.		Fixed seed=20260824, frozen manifest, replayable.			



Recalculation triggers: (1) official boundary released (2) official key-area polygons released (3) south gateway location (Issue #1029) clarified

On any of these, every column-A metric is recomputed and the figures, A3/A0 boards and HTML are regenerated.

Data: participant-supplied geometry/*.geojson, metrics.json, simulation.json. Readings come from offline synthetic rehearsal, not field measurement, and are not a performance commitment.